

# Contents

<b>1 Coalesce During Save Fix</b>	<b>1</b>
1.1 Issue . . . . .	1
1.2 Root Cause . . . . .	1
1.3 Solution . . . . .	2
1.3.1 Key Changes . . . . .	2
1.3.2 Implementation . . . . .	2
1.4 Files Modified . . . . .	3
1.4.1 1. skol_classifier/classifier_v2.py . . . . .	3
1.4.2 2. skol_classifier/classifier_v2.py . . . . .	3
1.5 Impact . . . . .	4
1.5.1 Before (Broken) . . . . .	4
1.5.2 After (Fixed) . . . . .	4
1.6 Workflow . . . . .	5
1.6.1 Complete Workflow with Coalescing . . . . .	5
1.7 Benefits . . . . .	6
1.8 Related Issues Fixed . . . . .	6
1.9 Testing . . . . .	6
1.10 Notes . . . . .	7

## 1 Coalesce During Save Fix

### 1.1 Issue

**Problem:** When coalesce\_labels=True, the coalescing was happening during predict(), which meant: 1. Users couldn't inspect individual line predictions 2. The notebook couldn't show predicted\_label for sample predictions 3. The output DataFrame had a completely different structure that broke inspection code

**User Report:** “In the same code we now get an error about missing predicted\_label. When rows are coalesced, they should all have the same predicted\_label and it should be preserved.”

### 1.2 Root Cause

The predict() method was calling \_format\_predictions(), which in turn called YeddaFormatter.format() with coalesce\_labels=True. This meant:

```
# In predict()
predictions_df = self._format_predictions(predictions_df)
# Returns coalesced format with no predicted_label column
```

When coalescing happens, the output structure changes from:

- **Line-level**: doc\_id, attachment\_name, value, line\_number, predicted\_label, annotated\_value

To: - **Coalesced**: doc\_id, attachment\_name, coalesced\_annotations (array of strings)

This broke notebook code that expected to inspect predicted\_label.

## 1.3 Solution

**Design Decision:** Keep coalescing as a **save-time operation** only, not a prediction-time operation.

### 1.3.1 Key Changes

1. **predict() always returns line-level data** - Even when coalesce\_labels=True, users can inspect individual predictions
2. **Coalescing happens only in save\_annotated()** - The output formatters apply coalescing when saving
3. **Better user experience** - Users can inspect predictions before saving, regardless of coalesce setting

### 1.3.2 Implementation

**Changed:** \_format\_as\_annotated() to never coalesce

```
# Before - coalescing happened during predict()
def _format_as_annotated(self, predictions_df: DataFrame) -> DataFrame:
    formatter = YeddaFormatter(
        coalesce_labels=self.coalesce_labels, # ✗ Applied coalescing
        line_level=self.line_level
    )
    return formatter.format(predictions_df)

# After - coalescing deferred to save time
def _format_as_annotated(self, predictions_df: DataFrame) -> DataFrame:
    """
    Format predictions as YEDDA-style annotated blocks.

    Note: This does NOT apply coalescing. Coalescing is only applied
    during save_annotated() to preserve line-level data for inspection.
    """
    formatter = YeddaFormatter(
        coalesce_labels=False, # ✓ Never coalesce in predict()
        line_level=self.line_level
```

```

        )
    return formatter.format(predictions_df)

Changed: _save_to_couchdb() to use CouchDBOutputWriter
# Before - used non-existent save_predictions method
def _save_to_couchdb(self, predictions: DataFrame) -> None:
    conn = CouchDBConnection(...)
    conn.save_predictions( # ❌ This method doesn't exist
        predictions,
        suffix=self.output_couchdb_suffix,
        coalesce_labels=self.coalesce_labels
    )

# After - uses proper output writer that handles coalescing
def _save_to_couchdb(self, predictions: DataFrame) -> None:
    from .output_formatters import CouchDBOutputWriter

    writer = CouchDBOutputWriter(
        couchdb_url=self.couchdb_url,
        database=self.couchdb_database,
        username=self.couchdb_username,
        password=self.couchdb_password
    )

    writer.save_annotated( # ✓ Applies coalescing here
        predictions,
        suffix=self.output_couchdb_suffix,
        coalesce_labels=self.coalesce_labels,
        line_level=self.line_level
    )

```

## 1.4 Files Modified

### 1.4.1 1. skol\_classifier/classifier\_v2.py

**Method:** \_format\_as\_annotated

**Changes:** - Line 565: Changed coalesce\_labels=self.coalesce\_labels to coalesce\_labels=False - Lines 559-562: Added docstring explaining that coalescing is deferred

### 1.4.2 2. skol\_classifier/classifier\_v2.py

**Method:** \_save\_to\_couchdb

**Changes:** - Replaced CouchDBConnection.save\_predictions()

(which doesn't exist) - With CouchDBOutputWriter.save\_annotated()  
(which handles coalescing properly) - Passes coalesce\_labels and  
line\_level to the writer

## 1.5 Impact

### 1.5.1 Before (Broken)

```
classifier = SkolClassifierV2(  
    spark=spark,  
    input_source='couchdb',  
    couchdb_database='mydb',  
    line_level=True,  
    coalesce_labels=True, # Coalescing during predict()  
    output_format='annotated'  
)  
  
predictions = classifier.predict()  
  
# ❌ BROKEN: No predicted_label column!  
predictions.select("doc_id", "attachment_name", "predicted_label").show()  
# Error: Column 'predicted_label' does not exist  
  
# Columns: doc_id, attachment_name, coalesced_annotations
```

### 1.5.2 After (Fixed)

```
classifier = SkolClassifierV2(  
    spark=spark,  
    input_source='couchdb',  
    couchdb_database='mydb',  
    line_level=True,  
    coalesce_labels=True, # Coalescing deferred to save()  
    output_format='annotated'  
)  
  
predictions = classifier.predict()  
  
# ✅ WORKS: predicted_label column exists!  
predictions.select("doc_id", "attachment_name", "predicted_label").show()  
# Shows individual line predictions  
  
# Columns: doc_id, attachment_name, value, line_number,  
#           predicted_label, annotated_value
```

```
# Coalescing happens here during save
classifier.save_annotated(predictions)
# Saved output has coalesced annotation blocks
```

## 1.6 Workflow

### 1.6.1 Complete Workflow with Coalescing

```
from pyspark.sql import SparkSession
from skol_classifier.classifier_v2 import SkolClassifierV2

spark = SparkSession.builder.getOrCreate()

# Configure with coalescing enabled
classifier = SkolClassifierV2(
    spark=spark,
    input_source='couchdb',
    couchdb_url='http://localhost:5984',
    couchdb_database='mydb',
    couchdb_username='admin',
    couchdb_password='password',
    couchdb_pattern='*.txt',
    line_level=True,
    coalesce_labels=True,           # Enable coalescing
    output_format='annotated',
    output_dest='couchdb',
    output_couchdb_suffix='.ann',
    model_type='logistic'
)

# 1. Train (if needed)
results = classifier.fit()

# 2. Predict - always returns line-level data
predictions = classifier.predict()

# 3. Inspect predictions - works because data is line-level
print("Sample predictions:")
predictions.select(
    "doc_id",
    "attachment_name",
    "predicted_label"
).show(5)

# 4. Show label distribution
predictions.groupBy("predicted_label").count().show()
```

```

# 5. Save - coalescing happens here
print("Saving with coalescing...")
classifier.save_annotated(predictions)
# Output in CouchDB has coalesced annotation blocks

```

## 1.7 Benefits

1. **Consistent API:** predict() always returns the same structure regardless of coalesce\_labels
2. **Inspectable Results:** Users can always examine individual predictions
3. **Flexible Output:** Users can inspect before saving, even with coalescing enabled
4. **Separation of Concerns:** Prediction and output formatting are separate operations
5. **Backward Compatible:** Existing code that doesn't use coalescing works the same way

## 1.8 Related Issues Fixed

This fix resolves:

- ✓ Missing predicted\_label column in predictions
- ✓ Inability to inspect predictions when coalescing enabled - ✓ Notebook code that expects line-level prediction data
- ✓ Incorrect use of non-existent CouchDBConnection.save\_predictions() method

## 1.9 Testing

To verify the fix:

```

from pyspark.sql import SparkSession
from skol_classifier.classifier_v2 import SkolClassifierV2

spark = SparkSession.builder.getOrCreate()

# Test with coalescing enabled
classifier = SkolClassifierV2(
    spark=spark,
    input_source='couchdb',
    couchdb_url='http://localhost:5984',
    couchdb_database='mydb',
    couchdb_pattern='*.txt',
    line_level=True,
    coalesce_labels=True, # Enable coalescing
    output_format='annotated',
)

```

```

        output_dest='couchdb',
        model_type='logistic'
    )

# Predict
predictions = classifier.predict()

# Test 1: Verify line-level structure
assert "predicted_label" in predictions.columns
assert "attachment_name" in predictions.columns
assert "line_number" in predictions.columns
print("✓ Line-level structure preserved")

# Test 2: Verify can inspect predictions
predictions.select("doc_id", "attachment_name", "predicted_label").show(5)
print("✓ Can inspect predictions")

# Test 3: Verify can analyze results
label_dist = predictions.groupBy("predicted_label").count()
label_dist.show()
print("✓ Can analyze prediction distribution")

# Test 4: Verify save works (coalescing applied here)
classifier.save_annotated(predictions)
print("✓ Save with coalescing successful")

```

## 1.10 Notes

- predict() returns line-level data even when coalesce\_labels=True
- Coalescing only happens during save\_annotated()
- Users can call predict() multiple times without re-coalescing
- The coalesce\_labels setting affects only the saved output format, not the prediction output