

# Contents

<b>1 CUDA Debugging Scripts</b>	<b>1</b>
1.1 Scripts . . . . .	1
1.1.1 1. test_cuda_init.py - Direct CUDA Driver API Test	1
1.1.2 2. test_tf_cuda_init.py - TensorFlow CUDA Test (Verbose)	2
1.2 Usage . . . . .	2
1.2.1 First Test: Raw CUDA	2
1.2.2 Second Test: TensorFlow	2
1.3 Interpreting Results . . . . .	2
1.3.1 Case 1: Both scripts succeed	2
1.3.2 Case 2: test_cuda_init.py succeeds, test_tf_cuda_init.py fails	3
1.3.3 Case 3: Both scripts fail	3
1.4 Verbose Logging Explained . . . . .	3
1.5 Saving Output . . . . .	3
1.6 What to Look For . . . . .	3
1.7 Next Steps . . . . .	4

## 1 CUDA Debugging Scripts

Two standalone scripts to help diagnose CUDA initialization issues.

### 1.1 Scripts

#### 1.1.1 1. test\_cuda\_init.py - Direct CUDA Driver API Test

Tests CUDA initialization using the raw CUDA driver API (via ctypes), bypassing TensorFlow entirely.

**Run:**

```
python3 test_cuda_init.py
```

**What it does:** - Loads libcuda.so.1 directly - Calls cuInit(0) to initialize CUDA - Reports success/failure with error codes - If successful, enumerates GPU devices and compute capabilities - If failed, shows nvidia-smi output and device permissions

**Exit codes:** - 0 - culinit() succeeded - 1 - culinit() failed or library not found

**This tests:** Whether CUDA driver itself can initialize, independent of TensorFlow.

### **1.1.2 2. test\_tf\_cuda\_init.py - TensorFlow CUDA Test (Verbose)**

Tests CUDA initialization through TensorFlow with maximum logging enabled.

#### **Run:**

```
python3 test_tf_cuda_init.py
```

**What it does:** - Sets maximum TensorFlow verbosity (TF\_CPP\_MIN\_LOG\_LEVEL=0)  
- Enables CUDA-specific module logging (TF\_CPP\_VMODULE=cuda\_diagnostics=10,...)  
- Imports TensorFlow and lists GPU devices - Attempts a simple GPU matrix multiplication - Provides diagnosis for common errors

**Exit codes:** - 0 - GPU operation succeeded - 1 - GPU initialization or operation failed

**This tests:** Whether TensorFlow can initialize and use CUDA/GPU.

## **1.2 Usage**

### **1.2.1 First Test: Raw CUDA**

Start with the direct CUDA test to see if the driver works at all:

```
cd /data/piggy/src/github.com/piggyatbaqaqi/skol
python3 test_cuda_init.py
```

**Expected for RTX 5090:** - ✓ Should succeed (culinit works, driver is fine) - Reports compute capability 12.0

### **1.2.2 Second Test: TensorFlow**

Then test TensorFlow specifically:

```
python3 test_tf_cuda_init.py 2>&1 | tee tf_cuda_test.log
```

The 2>&1 | tee captures all output (including stderr) to a file for review.

**Expected for RTX 5090:** - ✗ Will likely fail with CUDA\_ERROR\_INVALID\_PTX  
- Logs will show JIT compilation failure - Diagnosis will recommend CPU-only mode

## **1.3 Interpreting Results**

### **1.3.1 Case 1: Both scripts succeed**

✓ CUDA driver works ✓ TensorFlow can use GPU → **No issues**, GPU training will work

### 1.3.2 Case 2: test\_cuda\_init.py succeeds, test\_tf\_cuda\_init.py fails

- ✓ CUDA driver works ✗ TensorFlow can't use GPU → **Compute capability issue** (like RTX 5090 with TF 2.21) → Use CPU-only mode or wait for TensorFlow update

### 1.3.3 Case 3: Both scripts fail

✗ CUDA driver doesn't work ✗ TensorFlow can't use GPU → **Driver/system issue** → Check: - nvidia-smi works? - /dev/nvidia\* permissions? - Persistence mode: sudo nvidia-smi -pm 1 - Need reboot?

## 1.4 Verbose Logging Explained

The test\_tf\_cuda\_init.py script sets these environment variables:

- TF\_CPP\_MIN\_LOG\_LEVEL=0 - Show all log levels (INFO, WARNING, ERROR)
- TF\_CPP\_VMODULE=cuda\_diagnostics=10,... - Set module-specific verbosity to 10
- CUDA\_LAUNCH\_BLOCKING=1 - Make CUDA launches synchronous for clearer errors

This produces extensive output showing: - CUDA initialization steps - Driver/kernel version checks - Device enumeration - PTX/kernel compilation attempts - Exact point of failure

## 1.5 Saving Output

To save full verbose output for analysis:

```
# TensorFlow test with full logs
python3 test_tf_cuda_init.py 2>&1 | tee tf_verbose.log
```

```
# Direct CUDA test
python3 test_cuda_init.py 2>&1 | tee cuda_direct.log
```

Review the logs to find the exact error message and stack trace.

## 1.6 What to Look For

In test\_tf\_cuda\_init.py output, search for:

**Success indicators:** - "Created device /job:localhost/replica:0/task:0/device:GPU:0"  
- "/// GPU OPERATION SUCCESSFUL ///"

**Failure indicators:** - "CUDA\_ERROR\_INVALID\_PTX" → Compute capability too new - "CUDA\_ERROR\_INVALID\_HANDLE" → Usually follows INVALID\_PTX - "CUDA\_ERROR\_UNKNOWN" → Generic init failure - "failed call to cuInit" → Driver initialization failed - "JIT compiled from PTX" followed by errors → Compilation failed

## 1.7 Next Steps

After running both scripts:

1. **If test\_cuda\_init.py succeeds but test\_tf\_cuda\_init.py fails:**
  - Your hardware is fine
  - Use CPU-only mode: `os.environ['CUDA_VISIBLE_DEVICES'] = ''`
  - The auto-detection in rnn\_model.py should handle this
2. **If both fail:**
  - Check NVIDIA driver installation
  - Try: `sudo nvidia-smi -pm 1`
  - Reboot and test again
  - Check kernel module: `lsmod | grep nvidia`
3. **If both succeed:**
  - Something else is wrong with your training setup
  - The RTX 5090 should work (surprisingly!)
  - Check TensorFlow version and build info