

# Contents

<b>1 Extract Taxa Schema Type Mismatch Fixes</b>	<b>1</b>
1.1 Problem . . . . .	1
1.2 Root Causes . . . . .	1
1.2.1 What Taxon.as_row() Returns . . . . .	1
1.2.2 What the Schema Declared . . . . .	2
1.2.3 Why This Caused Failures . . . . .	2
1.3 The Fixes . . . . .	3
1.3.1 Fix 1: Integer Field Types (First Issue - 2763 Failures)	3
1.3.2 Fix 2: Allow None in Source Map (Second Issue - 5239 Failures)	3
1.3.3 Final Schema . . . . .	3
1.4 Testing . . . . .	3
1.5 Related Issue . . . . .	4
1.6 Files Modified . . . . .	4
1.7 Additional Notes . . . . .	4

## 1 Extract Taxa Schema Type Mismatch Fixes

### 1.1 Problem

The extract\_and\_save\_taxa\_pipeline() function was reporting failures (success=False) when trying to save taxa to CouchDB: - **First issue:** 2763 failures due to incorrect integer field types - **Second issue:** 5239 failures (after first fix) due to None values in source map

All extractions were failing, suggesting issues in the extraction stage, not the save stage.

### 1.2 Root Causes

**Multiple schema type mismatches** between the Spark DataFrame schema and the actual data types returned by Taxon.as\_row():

#### 1.2.1 What Taxon.as\_row() Returns

From taxon.py:60-85:

```
def as_row(self) -> Dict[str, None | str | int | Dict[str, None | str | int]]:  
    '''Convert this Taxon to a dictionary suitable for output.'''  
  
    # ... code ...
```

```

        retval: Dict[str, None | str | int | Dict[str, None | str | int]] = {
            'taxon': "\n".join((str(pp) for pp in self._nomenclatures)),
            'description': "\n".join((str(pp) for pp in self._descriptions)),
            'source': {
                'doc_id': source_doc_id,
                'url': source_url,
                'db_name': source_db_name,
            },
            'line_number': line_number,           # INTEGER (or None)
            'paragraph_number': pp.paragraph_number, # INTEGER (or None)
            'page_number': pp.page_number,         # INTEGER (or None)
            'empirical_page_number': pp.empirical_page_number, # INTEGER (or None)
        }
    return retval

```

The method returns: - **integers** (or None) for line\_number, paragraph\_number, page\_number - **string** (or None) for empirical\_page\_number (can be Roman numerals like "xvii") - **dict with None values** for source - specifically, source['url'] can be None

### 1.2.2 What the Schema Declared

From extract\_taxa\_to\_couchdb.py:290-298 (before fix):

```

extract_schema = StructType([
    StructField("taxon", StringType(), False),
    StructField("description", StringType(), False),
    StructField("source", MapType(StringType(), StringType()), False), # WRONG:
    StructField("line_number", StringType(), False),                  # WRONG: Should be I
    StructField("paragraph_number", StringType(), False),             # WRONG: Should be I
    StructField("page_number", StringType(), False),                  # WRONG: Should be I
    StructField("empirical_page_number", StringType(), True),         # Correct type, b
])

```

The schema had multiple issues: 1. **Integer fields declared as StringType**: The numeric fields were declared as strings 2. **MapType doesn't allow None values**: The source dict can have None for the url key, but MapType(StringType(), StringType()) doesn't allow null values 3. **Non-nullable integer fields**: The integer fields should be nullable since they can be None

### 1.2.3 Why This Caused Failures

When Spark tried to create the DataFrame at line 307:

```
taxa_df = spark.createDataFrame(taxa_rdd, extract_schema)
```

It encountered type mismatches and failed to create the DataFrame, causing all rows to fail. This failure occurred **before** any data reached the `save_taxa_to_couchdb_partition` function, which is why all 2763 records showed `success=False`.

## 1.3 The Fixes

### 1.3.1 Fix 1: Integer Field Types (First Issue - 2763 Failures)

Updated the schema to use `IntegerType` for numeric fields and make them nullable.

### 1.3.2 Fix 2: Allow None in Source Map (Second Issue - 5239 Failures)

Updated the `MapType` to allow `None` values using `valueContainsNull=True`.

### 1.3.3 Final Schema

```
from pyspark.sql.types import MapType, IntegerType

extract_schema = StructType([
    StructField("taxon", StringType(), False),
    StructField("description", StringType(), False),
    StructField("source", MapType(StringType(), StringType(), valueContainsNull=True),
    StructField("line_number", IntegerType(), True), # FIXED: IntegerType
    StructField("paragraph_number", IntegerType(), True), # FIXED: IntegerType
    StructField("page_number", IntegerType(), True), # FIXED: IntegerType
    StructField("empirical_page_number", StringType(), True), # Correct: StringType
])
```

Key changes: 1. **Import IntegerType:** Added to imports at line 288 2. **Changed integer field types:** Numeric fields now use `IntegerType` instead of `StringType` 3. **Made fields nullable:** Changed `nullable=False` to `nullable=True` for integer fields since they can be `None` 4. **Allow None in map values:** Added `valueContainsNull=True` to `MapType` for the source field, allowing `source['url']` to be `None` 5. **Keep empirical\_page\_number as StringType:** Since it can contain Roman numerals like "xvii"

## 1.4 Testing

After applying the fix, re-run the pipeline in the notebook:

```
taxa_df = extract_and_save_taxa_pipeline(
    spark=spark,
```

```

        ingest_couchdb_url=ingest_couchdb_url,
        ingest_db_name=ingest_db_name,
        taxon_couchdb_url=taxon_couchdb_url,
        taxon_db_name=taxon_db_name,
        ingest_username=ingest_username,
        ingest_password=ingest_password,
        pattern=pattern
    )

# Check for successes
successful = taxa_df.filter("success = true").count()
failed = taxa_df.filter("success = false").count()

print(f"Successful: {successful}")
print(f"Failed: {failed}")

# If there are still failures, check the error messages
if failed > 0:
    taxa_df.filter("success = false").select("error_message").distinct().show(tru

```

## 1.5 Related Issue

This same file also needs to be added to spark.submit.pyFiles in the notebook to avoid the ModuleNotFoundError: No module named 'extract\_taxa\_to\_couchdb' error when Spark workers try to deserialize UDFs.

Add this line to the Spark configuration in the notebook:

```
.config("spark.submit.pyFiles",
      f'{parent_path / "line.py"},{parent_path / "fileobj.py"},'
      f'{parent_path / "couchdb_file.py"},{parent_path / "finder.py"},'
      f'{parent_path / "taxon.py"},{parent_path / "paragraph.py"},'
      f'{parent_path / "label.py"},{parent_path / "file.py"},'
      f'{parent_path / "extract_taxa_to_couchdb.py"}' # ADD THIS LINE
)
```

## 1.6 Files Modified

- extract\_taxa\_to\_couchdb.py: Fixed schema definition

## 1.7 Additional Notes

- The type mismatch would have manifested as an error during DataFrame creation, not during the save operation
- All 2763 failures were due to this single schema issue

- The fix ensures the DataFrame schema matches the actual data structure returned by Taxon.as\_row()