

Contents

1 Model Persistence Fix for SkolClassifierV2	1
1.1 Issue	1
1.2 Solution	1
1.2.1 Key Changes	1
1.3 Implementation Details	2
1.3.1 Disk Save (classifier_v2.py:574-618)	2
1.3.2 Disk Load (classifier_v2.py:620-655)	2
1.3.3 Redis Save (classifier_v2.py:657-713)	3
1.3.4 Redis Load (classifier_v2.py:715-768)	3
1.4 Metadata Format	3
1.5 Benefits	4
1.6 Testing	4
1.7 Files Modified	5
1.8 Related Issues Fixed	5
1.9 Comparison with Original Approach	5
1.9.1 Before (Broken):	5
1.9.2 After (Fixed):	5
1.10 Notes	6

1 Model Persistence Fix for SkolClassifierV2

1.1 Issue

Error: TypeError: cannot pickle '_thread.RLock' object

Location: skol_classifier/classifier_v2.py:633

Cause: The code was trying to use Python's pickle module to serialize PySpark ML models directly. PySpark models contain Java objects (including thread locks) that cannot be pickled.

1.2 Solution

Changed from using pickle to using PySpark's native model persistence methods, matching the approach used in the original SkolClassifier.

1.2.1 Key Changes

1. **Disk Storage:** Use PySpark's `.save()` method to save models to a directory structure

2. **Redis Storage:** Create a tar.gz archive of the saved models and store in Redis
3. **Metadata:** Store configuration and label mappings as JSON alongside models

1.3 Implementation Details

1.3.1 Disk Save (classifier_v2.py:574-618)

```
def _save_model_to_disk(self) -> None:
    """Save model to disk using PySpark's native save."""
    # Create directory structure
    model_dir = model_path.parent / model_path.stem

    # Save feature pipeline using PySpark
    self._feature_pipeline.save(str(pipeline_path))

    # Save classifier model using PySpark
    classifier_model.save(str(classifier_path))

    # Save metadata as JSON
    json.dump(metadata, f, indent=2)
```

Directory Structure:

```
model_name/
└── feature_pipeline/      # PySpark pipeline model
└── classifier_model/     # PySpark classifier model
└── metadata.json          # Configuration and labels
```

1.3.2 Disk Load (classifier_v2.py:620-655)

```
def _load_model_from_disk(self) -> None:
    """Load model from disk using PySpark's native load."""
    # Load using PySpark
    self._feature_pipeline = PipelineModel.load(str(pipeline_path))
    classifier_model = PipelineModel.load(str(classifier_path))

    # Load metadata
    metadata = json.load(f)

    # Recreate SkolModel wrapper
    self._model = SkolModel(...)
    self._model.set_model(classifier_model)
```

1.3.3 Redis Save (classifier_v2.py:657-713)

```
def _save_model_to_redis(self) -> None:
    """Save model to Redis using tar archive."""
    # Save models to temp directory using PySpark
    self._feature_pipeline.save(str(pipeline_path))
    classifier_model.save(str(classifier_path))

    # Save metadata
    json.dump(metadata, f)

    # Create tar.gz archive
    with tarfile.open(fileobj=archive_buffer, mode='w:gz') as tar:
        tar.add(temp_path, arcname='.')

    # Save to Redis
    self.redis_client.set(self.redis_key, archive_data)
```

1.3.4 Redis Load (classifier_v2.py:715-768)

```
def _load_model_from_redis(self) -> None:
    """Load model from Redis tar archive."""
    # Get from Redis
    serialized = self.redis_client.get(self.redis_key)

    # Extract tar archive to temp directory
    with tarfile.open(fileobj=archive_buffer, mode='r:gz') as tar:
        tar.extractall(temp_path)

    # Load models using PySpark
    self._feature_pipeline = PipelineModel.load(str(pipeline_path))
    classifier_model = PipelineModel.load(str(classifier_path))

    # Load metadata and recreate wrapper
    self._model = SkolModel(...)
    self._model.set_model(classifier_model)
```

1.4 Metadata Format

The metadata.json file stores:

```
{
    "label_mapping": {
        "Label1": 0,
        "Label2": 1
    },
}
```

```

    "config": {
        "line_level": true,
        "use_suffixes": true,
        "min_doc_freq": 2,
        "model_type": "logistic",
        "model_params": {
            "maxIter": 10,
            "regParam": 0.01
        }
    },
    "version": "2.0"
}

```

1.5 Benefits

1. **Compatibility:** Uses PySpark's native serialization which properly handles Java objects
2. **Portability:** Saved models can be loaded on different machines with the same PySpark version
3. **Flexibility:** Models stored in Redis are in the same format, just compressed
4. **Consistency:** Matches the proven approach used in SkolClassifier V1

1.6 Testing

To test model persistence:

```

from pyspark.sql import SparkSession
from skol_classifier.classifier_v2 import SkolClassifierV2

spark = SparkSession.builder.getOrCreate()

# Train and save to disk
classifier = SkolClassifierV2(
    spark=spark,
    input_source='files',
    file_paths=['data/*.ann'],
    model_storage='disk',
    model_path='models/my_model.pkl', # Creates models/my_model/ directory
    line_level=True,
    use_suffixes=True,
    model_type='logistic'
)
results = classifier.fit() # Automatically saves to disk

```

```

# Load from disk
classifier2 = SkolClassifierV2(
    spark=spark,
    input_source='files',
    file_paths=['data/*.txt'],
    model_storage='disk',
    model_path='models/my_model.pkl',
    auto_load_model=True, # Automatically loads on init
    line_level=True
)
predictions = classifier2.predict()

```

1.7 Files Modified

1. skol_classifier/classifier_v2.py
 - `_save_model_to_disk()` - Lines 574-618
 - `_load_model_from_disk()` - Lines 620-655
 - `_save_model_to_redis()` - Lines 657-713
 - `_load_model_from_redis()` - Lines 715-768

1.8 Related Issues Fixed

This fix also resolves:

- ✓ Cannot serialize PySpark models - ✓
- Thread lock pickling errors - ✓
- Java object serialization issues - ✓
- Model portability across machines

1.9 Comparison with Original Approach

1.9.1 Before (Broken):

```

# ❌ Tried to pickle PySpark models directly
model_data = {
    'model': self._model, # Contains unpicklable Java objects
    'feature_pipeline': self._feature_pipeline
}
pickle.dump(model_data, f) # TypeError: cannot pickle '_thread.RLock'

```

1.9.2 After (Fixed):

```

# ✅ Use PySpark's native save
self._feature_pipeline.save(str(pipeline_path)) # Works!
classifier_model.save(str(classifier_path)) # Works!
json.dump(metadata, f) # Python-only data

```

1.10 Notes

- Model directories are created with the same base name as the `model_path` (e.g., `models/my_model.pkl` → `models/my_model/`)
- Temporary directories are properly cleaned up after Redis operations
- The `_feature_extractor` must be available when loading to determine the correct features column name