

# Contents

<b>1 PDF Section Extractor - YEDDA Annotation Support</b>	<b>1</b>
1.1 Overview . . . . .	1
1.2 YEDDA Format . . . . .	2
1.2.1 Features Supported . . . . .	2
1.3 Usage . . . . .	2
1.3.1 Basic Example . . . . .	2
1.3.2 Example Output . . . . .	2
1.4 Nested Annotations . . . . .	3
1.4.1 Example . . . . .	3
1.5 DataFrame Schema . . . . .	3
1.6 Working with Labels . . . . .	4
1.6.1 Filter by Label . . . . .	4
1.6.2 Label Statistics . . . . .	4
1.6.3 Combine with Section Names . . . . .	5
1.7 Implementation Details . . . . .	5
1.7.1 Parsing Algorithm . . . . .	5
1.7.2 Annotation Markers . . . . .	5
1.7.3 Label Lookup . . . . .	5
1.8 Examples . . . . .	5
1.8.1 Example 1: Training Data with Labels . . . . .	5
1.8.2 Example 2: Quality Control . . . . .	6
1.8.3 Example 3: Export Annotations . . . . .	6
1.9 Integration with SkolClassifierV2 . . . . .	6
1.10 Compatibility . . . . .	7
1.10.1 Backward Compatibility . . . . .	7
1.10.2 Migration . . . . .	7
1.11 Testing . . . . .	7
1.12 Limitations . . . . .	7
1.13 Future Enhancements . . . . .	8
1.14 See Also . . . . .	8

## 1 PDF Section Extractor - YEDDA Annotation Support

### 1.1 Overview

The PDFSectionExtractor now automatically parses YEDDA (Yet Another Entity Detection and Annotation) format annotations from PDF text and includes the active label for each section in the output DataFrame.

**Date:** 2025-12-22 **Breaking Changes:** None (backward compatible -

new nullable field)

## 1.2 YEDDA Format

YEDDA annotations use the format:

```
[@ text content  
#Label*]
```

### 1.2.1 Features Supported

- **Nested Annotations:** When annotations are nested, the inner-most label takes precedence
- **Multi-line Annotations:** Annotations can span multiple lines
- **Cross-Page Support:** Annotations can cross page boundaries within the same PDF file
- **File Boundary:** Annotations do NOT cross file boundaries (each PDF is processed independently)

## 1.3 Usage

### 1.3.1 Basic Example

```
from pyspark.sql import SparkSession  
from pdf_section_extractor import PDFSectionExtractor  
  
spark = SparkSession.builder.appName("PDFExtractor").getOrCreate()  
extractor = PDFSectionExtractor(spark=spark)  
  
# Extract from PDF with YEDDA annotations  
df = extractor.extract_from_document(  
    database='skol_dev',  
    doc_id='document-id'  
)  
  
# DataFrame now includes 'label' column  
df.select("value", "section_name", "label").show()
```

### 1.3.2 Example Output

Given a PDF with YEDDA annotations:

--- PDF Page 1 ---

```
[@ This is the introduction section.  
#Introduction*]
```

```
[@ Glomus mosseae Nicolson & Gerdemann, 1963.
```

```
#Nomenclature*]
```

```
[@ Spores formed singly or in clusters.
```

```
#Description*]
```

This is unannnotated conclusion text.

The resulting DataFrame:

value	line_number	section_name	label
This is the introduction section.	3	NULL	Introduction
Glomus mosseae Nicolson & Gerdemann, 1963.	6	NULL	Nomenclature
Spores formed singly or in clusters.	9	NULL	Description
This is unannotated conclusion text.	11	NULL	NULL

## 1.4 Nested Annotations

YEDDA annotations can nest, and the innermost label takes precedence:

```
[@ Outer annotation text  
[@ Inner nested annotation  
#InnerLabel*]  
Back to outer annotation  
#OuterLabel*]
```

Result: - Lines in inner annotation: label = "InnerLabel" - Lines in outer annotation only: label = "OuterLabel"

### 1.4.1 Example

```
Line 1: [@ This is nomenclature # OuterLabel = Nomenclature  
Line 2: [@ This is a description # InnerLabel = Description  
Line 3: of the species. # InnerLabel = Description  
Line 4: #Description*] # InnerLabel = Description  
Line 5: Back to nomenclature. # OuterLabel = Nomenclature  
Line 6: #Nomenclature*] # OuterLabel = Nomenclature
```

Labels assigned: - Lines 1: Nomenclature - Lines 2-4: Description (innermost wins!) - Lines 5-6: Nomenclature

## 1.5 DataFrame Schema

The DataFrame schema now includes a label field:

```

StructType([
    StructField("value", StringType(), False),
    StructField("doc_id", StringType(), False),
    StructField("attachment_name", StringType(), False),
    StructField("paragraph_number", IntegerType(), False),
    StructField("line_number", IntegerType(), False),
    StructField("page_number", IntegerType(), False),
    StructField("empirical_page_number", IntegerType(), True), # Nullable
    StructField("section_name", StringType(), True), # Nullable
    StructField("label", StringType(), True) # Nullable - YEDDA annotation
])

```

## 1.6 Working with Labels

### 1.6.1 Filter by Label

```

# Get all nomenclature sections
nomenclature_df = df.filter(df.label == "Nomenclature")

# Get all annotated sections
annotated_df = df.filter(df.label.isNotNull())

# Get unannotated sections
unannotated_df = df.filter(df.label.isNull())

```

### 1.6.2 Label Statistics

```

from pyspark.sql.functions import count

# Count sections by label
label_counts = df.groupBy("label").agg(count("*").alias("count"))
label_counts.show()

# Output:
# +-----+-----+
# |      label|count|
# +-----+-----+
# |      NULL|  150|
# | Nomenclature |   45|
# | Introduction|   12|
# | Description|   78|
# +-----+-----+

```

### 1.6.3 Combine with Section Names

```
# Sections with both labels and section names
df.filter(
    df.label.isNotNull() & df.section_name.isNotNull()
).select("value", "section_name", "label").show()
```

## 1.7 Implementation Details

### 1.7.1 Parsing Algorithm

1. **Line-by-Line Parsing:** Text is split into lines (1-indexed)
2. **Stack-Based Tracking:** Annotation markers are tracked using a stack
3. **Label Assignment:** When an annotation closes, all lines in its range get the label
4. **Nesting Resolution:** Innermost labels are assigned first (over-writing outer labels)

### 1.7.2 Annotation Markers

- **Start marker:** [@ - Pushes a new annotation onto the stack
- **End marker:** #Label\*] - Pops from stack and assigns label to all lines in range

### 1.7.3 Label Lookup

For each section/paragraph:

- The line\_number field indicates the first line of the section
- The label field contains the YEDDA label active at that line
- If no annotation is active, label is NULL

## 1.8 Examples

### 1.8.1 Example 1: Training Data with Labels

```
# Load annotated training data
training_df = extractor.extract_from_document(
    database='training_data',
    doc_id='annotated_article_001'
)

# Filter to only labeled data for training
labeled_df = training_df.filter(training_df.label.isNotNull())

# Train classifier
from skol_classifier.classifier_v2 import SkolClassifierV2
```

```

classifier = SkolClassifierV2(
    spark=spark,
    input_source='dataframe', # Hypothetical future feature
    use_suffixes=True,
    model_type='logistic'
)
classifier.fit(labeled_df.select("value", "label"))

```

### 1.8.2 Example 2: Quality Control

```

# Check annotation coverage
total = df.count()
labeled = df.filter(df.label.isNotNull()).count()
coverage = (labeled / total) * 100

print(f"Annotation coverage: {coverage:.1f}%")
print(f"Labeled sections: {labeled}/{total}")

```

### 1.8.3 Example 3: Export Annotations

```

# Export to JSON for review
annotations = df.filter(df.label.isNotNull()).select(
    "value", "label", "page_number", "line_number"
).collect()

import json
with open('annotations.json', 'w') as f:
    json.dump([row.asDict() for row in annotations], f, indent=2)

```

## 1.9 Integration with SkolClassifierV2

YEDDA labels can be used as training labels for the classifier:

```
from skol_classifier.classifier_v2 import SkolClassifierV2
```

```

# Extract with YEDDA annotations
classifier = SkolClassifierV2(
    spark=spark,
    input_source='couchdb',
    couchdb_url='http://localhost:5984',
    couchdb_database='training_docs',
    extraction_mode='section', # Use PDF section extraction
    use_suffixes=True,

```

```

        model_type='logistic'
    )

# Load sections with YEDDA labels
sections_df = classifier.load_raw()

# Filter to labeled sections for training
training_df = sections_df.filter(sections_df.label.isNotNull())

# Rename 'label' to match classifier expectations
training_df = training_df.withColumnRenamed("label", "label_from_yedda")

# Train on labeled data
# (Implementation depends on classifier API)

```

## 1.10 Compatibility

### 1.10.1 Backward Compatibility

- **Fully compatible:** Existing code continues to work
- **New field:** The label column is nullable and won't affect existing queries
- **No breaking changes:** All existing DataFrame operations work as before

### 1.10.2 Migration

No migration required. Existing code will see NULL values in the new label column for PDFs without YEDDA annotations.

## 1.11 Testing

Run the test suite:

```
python test_yedda_pdf_integration.py
```

Tests cover: - Basic annotation parsing - Nested annotation handling (innermost wins) - Multi-line annotations - DataFrame integration - NULL labels for unannotated text

## 1.12 Limitations

1. **Format:** Only supports YEDDA format [@ text #Label\*]
2. **File Boundaries:** Annotations cannot cross file boundaries
3. **Validation:** No validation that closing labels match opening markers

4. **Error Handling:** Malformed annotations are silently ignored

## 1.13 Future Enhancements

Possible improvements: 1. **Validation:** Warn about unclosed or mismatched annotations 2. **Alternative Formats:** Support other annotation formats (XML, JSON) 3. **Annotation Metadata:** Include annotation confidence or source 4. **Label Hierarchies:** Support hierarchical label structures 5. **Annotation Conflicts:** Handle overlapping annotations more explicitly

## 1.14 See Also

- yedda\_parser/README.md - YEDDA parser module
- PDF\_SECTION\_EXTRACTOR\_SUMMARY.md - Complete feature summary
- PDF\_PICTURE\_CAPTION\_EXTRACTION.md - Figure caption handling
- CLASSIFIER\_V2\_TOKENIZER\_UPDATE.md - Section mode classifier

---

**Status:** Complete and Tested **Version:** Added 2025-12-22

**Breaking Changes:** None **New Features:** - YEDDA annotation parsing

- Nested annotation support (innermost label wins) - Multi-line annotation support - Cross-page annotation support - Nullable label field in DataFrame schema - Full test coverage

**Known Limitations:** - Annotations cannot cross file boundaries - No validation of annotation format - Malformed annotations silently ignored